

High Throughput Computing with HTCondor

Presented by:

Sabbir Ahmed Sibli (20266027)

Nuray Jannat (20266028)

Hasibul Islam Sifat (20266008)

Taniya Sultana Jabin (20266016)

Sahiba Tasneem (20266022)


Table of Contents

- ❏ Intro: High Throughput Computing (HTC)
- ❏ Intro: HTCondor
- ❏ HTCondor Terminologies
- ❏ Steps of Applying HTCondor
- ❏ Managing Workflow: DAGMan Tool
- ❏ Exceptional Features of HTCondor
- ❏ References

Intro: High Throughput Computing (HTC)

- ❑ HPC deals with one particular large problem.
- ❑ HTC deals with problem which can be broken into many smaller independent problems.
- ❑ Specifically, HTC means getting lots of work done in a specific time unit.
- ❑ HTC distributes tasks over many computers (e.g., idle desktop computers, dedicated servers, or cloud-based resources).

Intro: High Throughput Computing (HTC) contd.



High-throughput computing resources speed up the process of knee modeling simulations by sending jobs to many computers rather than just one. (Credit: University of Wisconsin-Madison)

Intro: HTCondor

- ❑ HTCondor "has enabled ordinary users to do extraordinary computing"
- ❑ It is a Distributed-batch computing system
- ❑ Users submit a job to HTCondor, then HTCondor chooses when and where to run the job based upon the job requirement and available worker machines.
- ❑ It monitors the jobs' progress, and notifies the user upon completion.
- ❑ The development of HTCondor started in 1998 by University of Wisconsin-Madison.

Intro: HTCondor contd.

HTCondor – an open-source workload management software for HTC jobs
(developed at the [University of Wisconsin – Madison](#))




Image Credit: fakecineaste blog

HTCondor Terminologies

- ❑ **Job**
 - ❑ The Large task a user wants HTCondor to run.
- ❑ **Batch/Shell Scripts**
 - ❑ A file with instructions for a computer to execute.
- ❑ **Submit Machine**
 - ❑ One or more computers used by HTCondor to send jobs out at the start of a run and collect jobs at the end of a run.
- ❑ **Class Ads**
 - ❑ A language used by HTCondor to advertise resources on machines running HTCondor.
- ❑ **Submit File**
 - ❑ A file submitted by the user to the central manager. This file tells HTCondor the job requirements

Steps of Applying HTCondor


- ❑ Define the computing problem.
 - ❑ Choosing appropriate problems which suit to run with HTC rather than HPC.
- ❑ Discretize the problem into smaller jobs that may be run on HTCondor.
 - ❑ Breaking the large single task into many independent task that they can run independently.
- ❑ Process the Inputs
 - ❑ After discretizing, having a way to identify files by some naming convention.
- ❑ Run the job on HTCondor
 - ❑ Placement of a submit file and corresponding files on the central manager.
- ❑ Post process the Data.
 - ❑ This step consists of combining the output from the previously run HTCondor jobs. Often, this step is most easily done using a scripting language such as Python or R.

Managing Workflow: DAGMan Tool


- ❑ It is a HTCondor tool that allows multiple jobs to be organized in **workflows**.
- ❑ It is represented as a directed acyclic graph (DAG).
- ❑ A DAGMan workflow automatically submits jobs in a particular order, such that certain jobs need to complete before others start running.
- ❑ This allows the outputs of some jobs to be used as inputs for others, and makes it easy to replicate a workflow multiple times in the future.

Managing Workflow: DAGMan Tool contd.

- ❑ A DAGMan workflow is described in a **DAG input file**.
- ❑ The input file specifies the nodes of the DAG. A **node** within a DAG represents a unit of work. It contains the following:
 - **Job**: An HTCondor job, defined in a submit file.
 - **Pre-Script**: A script that runs before the job starts. Typically used to verify that all inputs are valid.
 - **Post-Script**: A script that runs after the job finishes. Typically used to verify outputs and clean up temporary files.



Managing Workflow: DAGMan Tool contd.



A simple diamond-shaped DAG. **Edge** in DAGMan describes a dependency between two nodes (i.e. parent, child).

A very simple DAG input file for this diamond-shaped DAG is:

```
# File name: diamond.dag
```

```
JOB A A.condor
```

```
JOB B B.condor
```

```
JOB C C.condor
```

```
JOB D D.condor
```

```
PARENT A CHILD B C
```

```
PARENT B C CHILD D
```

Exceptional Features of HTCondor

❑ Scalability

- ❑ An HTCondor pool is horizontally scalable to hundreds of thousands of execute cores and a similar number of jobs.

❑ Security

- ❑ HTCondor can be configured to use strong authentication and encryption between the services on remote machines used to manage jobs.

❑ No Special Program required to run

- ❑ No special programming is required to use HTCondor.

❑ Flexible Policy Mechanisms

- ❑ HTCondor allows users to specify very flexible policies for how they want jobs to be run.

“The purpose of computing is insight, not numbers.”

- Richard Hamming

References

- 1) Erickson, R.A., Fienen, M.N., McCalla, S.G., Weiser, E.L., Bower, M.L., Knudson, J.M. and Thain, G., 2018. Wrangling distributed computing for high-throughput environmental science: An introduction to HTCondor. PLoS computational biology, 14(10), p.e1006468.
- 2) <https://htcondor.readthedocs.io/>
- 3) Thain D, Tannenbaum T, Livny M. Distributed computing in practice: the Condor experience. Concurrency and Computation: Practice and Experience. 2005; 17:323–356.
- 4) <http://fakecineaste.blogspot.com/2019/06/high-throughput-computing-htc.html>

Thanks!

